The Universal Deviation Principle: Multi-Spectrum Sensitivity Guarantees for Anomaly Detection

Odeyemi Olusegun Israel
Independent Researcher
segettii@gmail.com
ORCID: 0009-0004-2309-6436

February 26, 2026

Abstract

We introduce the **Universal Deviation Principle (UDP)**, a theoretically grounded framework for multi-perspective anomaly detection. The core principle is that, under explicit assumptions on operator-family richness and anomaly-normal separation in feature space, a point which is *meaningfully different from the normal set* cannot simultaneously match normality under all structural perspectives. We formalise this via *deviation-separating operator families*, collections of differentiable maps $\{\Phi_k\}_{k=1}^K$ whose stacked Jacobian has full column rank, yielding local distinguishability around a reference normal state $\boldsymbol{\mu}$. To connect the sensitivity theory to detection, we define a normal set $\mathcal{S}$ and a deviation score based on distance of $\Phi(\mathbf{x})$ to $\Phi(\mathcal{S})$; under an explicit separation condition between normal and anomaly sets in operator space, at least one operator deviates by a positive margin. All theoretical results are conditional on three stated assumptions (observability, operator-family richness, calibration discipline). We prove five results: (1) local distinguishability via the Inverse Function Theorem; (2) global sensitivity via compactness and the Extreme Value Theorem with a uniform bound $\underline{\sigma} > 0$; (3) finite-sample guarantees via Hoeffding concentration; (4) fusion weights maximising worst-case directional sensitivity via a semidefinite programme; and (5) a greedy operator selection algorithm with $(1 - 1/e)$-approximation guarantee via monotone submodularity. We instantiate the framework with five spectrum projection operators (statistical, dynamical, spectral, geometric, exponential) and introduce a **Boundary-Centred Dimension Magnifier** paired with QDA that, given labels, achieves strong supervised performance. The unsupervised MDN score provides a label-free alternative (mean AUC 0.829 across five benchmarks).

We also present the **GravityEngine**, an N-body scoring algorithm whose force potentials we derive in closed form. The attraction potential involves the error function $\mathrm{erf}(r/\sigma)$. With backtracking line-search, the energy decreases monotonically at every observed iteration; the descent guarantee holds conditional on the line-search finding a step within the halving budget. Empirically, UDL with experimental operator stacks achieves mean AUC 0.960 across five benchmarks (two synthetic, three ODDS datasets), ranking third among all tested methods—within 0.007 of k-NN—under a single hyperparameter set with zero learnable parameters. The supervised UDL-QDA-Magnified variant reaches 0.895 mean AUC on the same five benchmarks. GravityEngine reaches AUC 0.899 on mammography with empirically verified monotone energy decrease (14/14 iteration steps). We numerically verify the theorem conditions on three independent ODDS benchmarks via Monte Carlo checks and observe no violations in our samples. Code is available at https://github.com/segetii/fintech.

1 Introduction

Every anomaly detector makes assumptions about what anomalies look like. Isolation Forest [Liu et al., 2008] assumes geometric isolation. LOF [Breunig et al., 2000] assumes local density deviation. Deep SVDD [Ruff et al., 2018] assumes separation from a learned hypersphere. Each works well when its assumptions hold, and each has blind spots when they do not. An anomaly that mimics normal local density will evade LOF entirely, yet the same anomaly may be obvious under spectral analysis or geometric projection. This is not a minor inconvenience; it is a structural limitation that recurs across fraud detection [Ahmed et al., 2016], medical diagnosis [Fernando et al., 2021], cybersecurity [Chandola et al., 2009], industrial monitoring [Pang et al., 2021], and aerospace telemetry [Hundman et al., 2018].

A natural question follows: can we state conditions under which combining multiple mathematical perspectives makes anomalies detectable? This paper answers in the affirmative under explicit assumptions on operator-family richness, feature observability, and calibration discipline. We introduce the **Universal Deviation Principle (UDP)**, which provides local, global, finite-sample, and algorithmic guarantees for operator-family sensitivity, and a set-based separation condition that turns sensitivity into a detection statement. All theoretical results are conditional on these assumptions; we state them explicitly in Section 3.

Contributions. Our contributions are:

1. **Theoretical framework** (Section 3): Three explicit assumptions (observability, operator-family richness, calibration discipline) under which five conditional results establish *sensitivity/distinguishability* for deviation-separating operator families, plus an explicit set-based separation condition that yields a detection guarantee. This includes local (Theorem 6) and global (Theorem 7) sensitivity bounds, finite-sample bounds (Theorem 8), fusion weights maximising worst-case sensitivity via SDP (Theorem 10), and greedy selection with $(1 - 1/e)$-approximation (Theorem 12). The individual proof techniques (IFT, EVT, Hoeffding, SDP relaxation, submodular optimisation) are classical; the contribution is the *framework* that threads them into end-to-end conditions for detectability.
2. **GravityEngine** (Section 4): An N-body scoring algorithm with a closed-form energy functional (attraction via erf potential, logarithmic repulsion, quadratic confinement). Backtracking line-search yields monotone energy descent at every observed iteration (Theorem 13); we state the descent guarantee conditional on the line-search finding a step within the halving budget.
3. **Multi-spectrum architecture** (Section 5): Five composable spectrum operators, a Boundary-Centred Dimension Magnifier with saturating $\tanh$ non-linearity, and QDA-Magnified projection achieving state-of-the-art AUC.
4. **Empirical evaluation** (Section 6): Five datasets across four domains, a single hyperparameter set throughout, comparisons against seven baselines including Deep SVDD and ECOD, 5-fold cross-validation on mammography, per-law decomposition analysis (including pendigits), and numerical checks of theorem conditions on three ODDS datasets.

2 Related Work

Classical methods. Distance-based approaches (k-NN [Ramaswamy et al., 2000], Mahalanobis) and density-based methods (LOF [Breunig et al., 2000], DBSCAN [Ester et al., 1996]) operate in geometric space. Tree-based methods like Isolation Forest [Liu et al., 2008] exploit random partitioning. All of these impose a single mathematical viewpoint on the data.

Deep approaches. Autoencoder-based detectors [An and Cho, 2015, Zhou and Paffenroth, 2017] flag high reconstruction error. Deep SVDD [Ruff et al., 2018] learns a minimum-volume hypersphere. GANs [Schlegl et al., 2017] model the normal distribution and flag deviations from it. These capture complex patterns, but they still produce a single representation of normality.

Modern unsupervised detectors. ECOD [Li et al., 2022] uses empirical cumulative distribution functions per feature. COPOD [Li et al., 2020] extends this with copula-based dependence modelling. Both achieve competitive performance in $O(n)$ time. Neither provides theoretical detectability guarantees.

OOD detection. Hendrycks and Gimpel [Hendrycks and Gimpel, 2017] showed that softmax probabilities give a usable OOD baseline; ODIN [Liang et al., 2018] and energy-based methods [Liu et al., 2020] improved on it. These produce a single OOD score. UDL differs because it decomposes *which mathematical aspect* makes an observation anomalous.

Ensemble and multi-view methods. Feature bagging [Lazarevic and Kumar, 2005] and detector aggregation [Aggarwal, 2015] combine multiple instances of similar detectors. Krogh and Vedelsby [Krogh and Vedelsby, 1995]

proved that ensemble error decreases with member diversity. UDL's diversity operates at the *representation* level rather than the *model* level: it fuses mathematically distinct projection operators, not multiple copies of the same detector.

Physics-inspired clustering. Particle swarm optimisation [Kennedy and Eberhart, 1995], gravitational search [Rashedi et al., 2009], and charged system search [Kaveh and Talatahari, 2010] use physical analogies for optimisation. GravityEngine differs in two ways: its potential functions are derived from the force field in closed form with force-energy consistency ($F = -\nabla E$) verified analytically, and backtracking line-search enforces monotone energy descent at every observed iteration.

Benchmarking. PyOD [Zhao et al., 2019] implements over 30 detectors. ADBench [Han et al., 2022] provides 57 benchmark datasets. We evaluate on ODDS benchmarks [Rayana, 2016] and follow ADBench protocols.

3 Universal Deviation Principle: Theory

3.1 Scope and Assumptions

The theoretical results in this section are conditional on three explicit assumptions. We state them upfront so that the scope of each theorem is unambiguous.

A1 (Observability) Anomalies are detectable from available features: the anomaly definition is *feature-realisable*, meaning there exists a measurable difference between normal and anomalous observations in $\mathbb{R}^d$. This excludes "label-only" anomalies where $\mathbf{x} \in \mathrm{supp}(P_{\text{normal}})$ but the label is adverse for reasons invisible in the feature space.

A2 (Operator-family richness) For any $\mathbf{x}$ that violates a normal model (see A1), there exists at least one operator Φ_k in the family for which the induced deviation exceeds a positive margin. Formally, the stacked Jacobian has full column rank (Theorem 1), and the separation condition of Theorem 4 holds with $\Delta > 0$. This assumption can fail if operators are redundant or insensitive on certain subspaces; we verify it empirically via Monte Carlo checks (Table 7).

A3 (Calibration discipline) All operator statistics, thresholds, and fusion weights are estimated on normal training data only. No test-set information (labels, statistics, or thresholds) is used in any step of score computation, operator selection, or decision boundary calibration.

What these assumptions exclude. A1 excludes adversarial label anomalies where $\mathbf{x}$ is indistinguishable from normals in feature space. A2 excludes degenerate operator families (e.g., K copies of the same linear projection). A3 excludes any form of test-set leakage. When all three hold, the theorems below provide conditional guarantees. When they do not, the framework reduces to an empirical multi-view anomaly detector with no formal guarantees—still useful, but not "proved."

3.2 Deviation-Separating Operator Families

Definition 1 (Deviation-Separating Family)**.** *Let $\{\Phi_k\}_{k=1}^K$, $\Phi_k : \mathbb{R}^d \to \mathbb{R}^{d_k}$, be continuously differentiable operators. The family is* deviation-separating *at $\boldsymbol{\mu} \in \mathbb{R}^d$ if:*

$$\bigcap_{k=1}^{K} \ker(D\Phi_k(\boldsymbol{\mu})) = \{\mathbf{0}\} \tag{1}$$

Equivalently (Theorem 2), the stacked Jacobian $J(\boldsymbol{\mu}) = [D\Phi_1(\boldsymbol{\mu})^\top, \ldots, D\Phi_K(\boldsymbol{\mu})^\top]^\top \in \mathbb{R}^{D\times d}$ has full column rank, where $D = \sum_k d_k$.

Lemma 2 (Rank Characterisation)**.** *$\{\Phi_k\}$ is deviation-separating at $\boldsymbol{\mu}$ if and only if $\mathrm{rank}(J(\boldsymbol{\mu})) = d$.*

The proof is immediate: $\ker(J(\boldsymbol{\mu})) = \bigcap_k \ker(D\Phi_k(\boldsymbol{\mu}))$.

3.3 From Sensitivity to Detection: A Normal Set $\mathcal{S}$

Theorems 6–8 are sensitivity statements around a reference point $\boldsymbol{\mu}$: they guarantee that *any* perturbation away from $\boldsymbol{\mu}$ produces a response in at least one operator under rank conditions. To interpret this as anomaly detection, we must specify what "normal" means (a set or distribution) and a detection statistic.

Definition 3 (Normal Set and Deviation Score)**.** *Let $\mathcal{S} \subset \mathbb{R}^d$ denote the set of normal states (e.g., the support of a normal distribution $P_{\mathcal{S}}$ or an empirical normal training set). Define the joint operator map $\Phi(\mathbf{x}) := (\Phi_1(\mathbf{x}), \ldots, \Phi_K(\mathbf{x})) \in \mathbb{R}^D$, $D = \sum_k d_k$. Define the set-distance deviation score*

$$s_{\mathcal{S}}(\mathbf{x}) := \mathrm{dist}(\Phi(\mathbf{x}), \Phi(\mathcal{S})) = \inf_{\mathbf{z}\in\mathcal{S}} \|\Phi(\mathbf{x}) - \Phi(\mathbf{z})\|_2 . \tag{2}$$

For per-operator analysis, define $s_{\mathcal{S},k}(\mathbf{x}) := \mathrm{dist}(\Phi_k(\mathbf{x}), \Phi_k(\mathcal{S}))$.

Theorem 4 (Separation in Operator Space Implies Detectability)**.** *Assume* $\mathcal{S}$ *is compact and each* Φ_k *is continuous. Let* $\mathcal{A} \subset \mathbb{R}^d$ *be a compact set of anomalies with* $\mathcal{A} \cap \mathcal{S} = \emptyset$*, and assume there exists a margin* $\Delta > 0$ *such that*

$$\mathrm{dist}(\Phi(\mathcal{A}), \Phi(\mathcal{S})) \geq \Delta. \tag{3}$$

Then for all $\mathbf{x} \in \mathcal{A}$*,* $s_{\mathcal{S}}(\mathbf{x}) \geq \Delta$ *and in particular there exists* $k \in \{1, \ldots, K\}$ *such that*

$$s_{\mathcal{S},k}(\mathbf{x}) \geq \Delta/\sqrt{K}. \tag{4}$$

Remark 5 (What the guarantee does and does not claim)**.** *The separation condition* (3) *is an explicit modelling assumption: if anomalies are not separated from normals in operator space (e.g., an adversary constructs* $\mathbf{x} \notin \mathcal{S}$ *with* $\Phi(\mathbf{x}) \in \Phi(\mathcal{S})$*), no operator-based detector can succeed. UDL provides (i) conditions under which perturbations are not invisible (rank conditions), and (ii) a clean separation-to-detection statement once a normal set* $\mathcal{S}$ *and a detection score are fixed.*

Empirical evidence for operator-space separation. To verify that the separation assumption $\Delta > 0$ holds in practice, we compute Cohen's d (standardised centroid distance) between anomaly and normal embeddings $\Phi(\mathcal{A})$, $\Phi(\mathcal{S})$ on all five benchmarks. The empirical values are: Synthetic $d = 14.6$, Mimic $d = 13.7$, Mammography $d = 3.3$, Shuttle $d = 0.43$, Pendigits $d = 0.81$. All are strictly positive, confirming that the joint operator map induces measurable separation on every dataset. The variation across datasets reflects detection difficulty: large d (synthetic, mimic) corresponds to trivial detection (AUC ≈ 1.0), while smaller d (shuttle, pendigits) correlates with harder tasks. Per-law analysis shows the geometric spectrum contributes the largest separation on four of five datasets, consistent with the per-law AUC results in Table 5.

3.4 Main Theorems

Theorem 6 (Universal Deviation, Local Form)**.** *Let* $\{\Phi_k\}_{k=1}^K$ *be a deviation-separating family at* $\boldsymbol{\mu}$ *with each* $\Phi_k \in C^1$*. Then there exists an open neighbourhood* $U \ni \boldsymbol{\mu}$ *such that for every* $\mathbf{x} \in U \setminus \{\boldsymbol{\mu}\}$*:*

$$\exists\, k \in \{1, \ldots, K\} : \quad \Phi_k(\mathbf{x}) \neq \Phi_k(\boldsymbol{\mu}) \tag{5}$$

Since $J(\boldsymbol{\mu})$ *has full column rank and singular values depend continuously on the matrix, there exists* $\varepsilon > 0$ *such that* $\mathrm{rank}(J(\mathbf{x})) = d$ *for all* $\mathbf{x} \in B(\boldsymbol{\mu}, \varepsilon)$*. This neighbourhood is precisely the* U *above.*

With $\sigma_{\min} := \sigma_{\min}(J(\boldsymbol{\mu})) > 0$*, a quantitative bound holds:*

$$\max_k \|\Phi_k(\mathbf{x}) - \Phi_k(\boldsymbol{\mu})\| \geq \frac{\sigma_{\min}}{\sqrt{K}} \|\mathbf{x} - \boldsymbol{\mu}\| - O(\|\mathbf{x} - \boldsymbol{\mu}\|^2) \tag{6}$$

Theorem 7 (Universal Deviation, Global Form)**.** *Let each* $\Phi_k \in C^2$*, and let* $\mathcal{K} \subset \mathbb{R}^d$ *be compact with* $\mathrm{rank}(J(\mathbf{x})) = d$ *for all* $\mathbf{x} \in \mathcal{K}$*. Define* $\underline{\sigma} := \min_{\mathbf{x} \in \mathcal{K}} \sigma_{\min}(J(\mathbf{x}))$*. Then* $\underline{\sigma} > 0$*, and for all* $\mathbf{x} \in \mathcal{K}$*:*

$$\max_k \|\Phi_k(\mathbf{x}) - \Phi_k(\boldsymbol{\mu})\| \geq \frac{\underline{\sigma}}{\sqrt{K}} \|\mathbf{x} - \boldsymbol{\mu}\| - C \|\mathbf{x} - \boldsymbol{\mu}\|^2 \tag{7}$$

where $C = \frac{1}{2} \sup_{\mathbf{z} \in \mathcal{K}} \left\| D^2 \Phi(\mathbf{z}) \right\|$ *is finite by continuity of* $D^2\Phi$ *on the compact set.*

Theorem 8 (Finite-Sample Deviation Guarantee)**.** *Let* $\hat{\boldsymbol{\mu}}_n = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$ *with* $\mathbf{x}_i \overset{iid}{\sim} P_{\mathcal{S}}$ *sub-Gaussian, population mean* $\boldsymbol{\mu}$*, covariance* Σ*. Let each* Φ_k *be* L_k*-Lipschitz,* $L = \max_k L_k$*. Then with probability* $\geq 1 - \delta$*:*

$$\max_k \|\Phi_k(\mathbf{x}) - \Phi_k(\hat{\boldsymbol{\mu}}_n)\| \geq \frac{\underline{\sigma}}{\sqrt{K}} \|\mathbf{x} - \boldsymbol{\mu}\| - \frac{L\sqrt{\mathrm{tr}(\Sigma)}\sqrt{2\ln(2d/\delta)}}{\sqrt{n}} \tag{8}$$

Corollary 9 (Sample Complexity)**.** *The bound* (8) *is non-trivial when* $n > 2L^2 K \cdot \mathrm{tr}(\Sigma) \cdot \ln(2d/\delta)/(\underline{\sigma}^2 \|\mathbf{x} - \boldsymbol{\mu}\|^2)$.

Theorem 10 (Worst-Case Sensitivity Maximisation via SDP)**.** *Define per-operator information matrices* $\mathcal{I}_k := D\Phi_k(\boldsymbol{\mu})^\top D\Phi_k(\boldsymbol{\mu}) \succeq 0$*. The worst-case detection sensitivity under weighted fusion* $S_{\mathbf{w}}[\mathbf{x}] = \sum_k w_k \|\Phi_k(\mathbf{x}) - \Phi_k(\boldsymbol{\mu})\|^2$ *is:*

$$\eta(\mathbf{w}) := \lambda_{\min}(\textstyle\sum_k w_k \mathcal{I}_k) \tag{9}$$

which is concave on the simplex Δ_K *(since* $\lambda_{\min}$ *is an infimum of linear functions). Geometrically,* $\eta(\mathbf{w})$ *is the worst-case directional sensitivity: the minimum squared response of the fused operator to any unit perturbation. Maximising* η *over* $\mathbf{w} \in \Delta_K$ *yields the weight vector* $\mathbf{w}^*$ *that makes the weakest detection direction as strong as possible, solvable as a standard SDP.*

Remark 11 (Finite-sample caveat for SDP weights)**.** *The SDP uses* $\mathcal{I}_k$ *matrices computed from* $D\Phi_k(\boldsymbol{\mu})$*, which in practice is estimated from data via finite differences. The "optimality" of* $\mathbf{w}^*$ *is therefore only as good as the Jacobian estimates; under finite-sample estimation noise, the true worst-case sensitivity may differ. We treat the SDP as a principled optimisation heuristic with population-level optimality, not a finite-sample guarantee, and verify empirically that SDP weights outperform equal weights on two of three datasets (Table 7).*

Theorem 12 (Greedy Operator Selection)**.** *Given* M *candidate operators with* $\mathcal{I}_k \succeq 0$ *for each* k *and budget* K*, selecting* S^* *to maximise* $\lambda_{\min}(\sum_{k\in S} \mathcal{I}_k)$ *is NP-hard. The surrogate* $f(S) = \log\det(\sum_{k\in S} \mathcal{I}_k)$ *is monotone submodular (monotonicity and diminishing returns both require* $\mathcal{I}_k \succeq 0$*), and greedy selection achieves:*

$$f(S_{\textit{greedy}}) \geq (1 - 1/e) \cdot f(S^*) \tag{10}$$

Surrogate gap. The $(1 - 1/e)$ guarantee holds for $f(S) = \log\det(\cdot)$, not directly for $\lambda_{\min}(\cdot)$. In practice the two objectives are correlated: maximising the log-determinant tends to raise the smallest eigenvalue because $\log\det = \sum \log \lambda_i$ penalises small eigenvalues logarithmically. Bounding the gap between the two objectives in the worst case remains open; we report both SDP-optimal and greedy performance empirically (Table 7).

All proofs are in Section A.

3.5 How $J(\mathbf{x})$ is computed in practice

For general (nonlinear) operators where analytic derivatives are inconvenient, we estimate each Jacobian $D\Phi_k(\boldsymbol{\mu})$ by central finite differences. Concretely, for input dimension d and reference point $\boldsymbol{\mu}$, we perturb coordinate j by $\varepsilon_j = \varepsilon(|\mu_j| + 10^{-6})$ and compute

$$\frac{\partial \Phi_k}{\partial x_j}(\boldsymbol{\mu}) \approx \frac{\Phi_k(\boldsymbol{\mu} + \varepsilon_j \mathbf{e}_j) - \Phi_k(\boldsymbol{\mu} - \varepsilon_j \mathbf{e}_j)}{2\varepsilon_j}. \tag{11}$$

We then stack the per-operator Jacobians vertically to form $J(\boldsymbol{\mu})$ and compute $\sigma_{\min}(J)$ via SVD. This matches the implementation used in our operator-diversity diagnostics (see the finite-difference Jacobian routine in `udl/energy.py`).

4 GravityEngine: N-Body Variational Scoring

The theoretical framework of Section 3 provides sensitivity/distinguishability guarantees conditional on Assumptions A1–A3 and, under the explicit separation condition of Theorem 4, a detectability statement. We now turn to the question of *how to score* observations. The GravityEngine implements the deviation principle through energy-inspired N-body dynamics.

4.1 Force Field

Each data point $\mathbf{x}_i$ is treated as a particle subject to three forces:

Radial anchoring. A spring-like restoring force toward the data centroid:

$$\mathbf{F}_{\text{radial}}(\mathbf{x}_i) = -\alpha(\mathbf{x}_i - \boldsymbol{\mu}) \tag{12}$$

Pairwise interaction. Short-range attraction with close-range repulsion between particles:

$$\mathbf{F}_{\text{pair}}(\mathbf{x}_i) = -\gamma \sum_{j \neq i} \left[\exp\left(-\frac{r_{ij}^2}{\sigma^2} \right) - \frac{\lambda}{r_{ij}} \right] \hat{\mathbf{r}}_{ij} \tag{13}$$

where $r_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|$ and $\hat{\mathbf{r}}_{ij}$ is the unit vector from j to i.

Deviation penalty. Connects to the UDL framework via the quadratic operator $\Phi(\mathbf{x}) = \mathbf{x}^2$. The deviation energy is $V_{\text{dev}} = \frac{\beta}{2}\left\|\mathbf{x}^2 - \boldsymbol{\mu}^2\right\|^2$; the force is its negative gradient:

$$F_{\text{dev},j}(\mathbf{x}_i) = -\frac{\partial V_{\text{dev}}}{\partial x_{ij}} = -2\beta\, x_{ij}(x_{ij}^2 - \mu_j^2) \tag{14}$$

i.e. $\mathbf{F}_{\text{dev}} = -2\beta\, \mathbf{x}_i \odot (\mathbf{x}_i^2 - \boldsymbol{\mu}^2)$, where $\odot$ is the Hadamard product.

After iterative evolution, normal points cluster together. Anomalies do not. They end up in isolated positions. The anomaly score is simply the displacement $\left\|\mathbf{x}_i^{(\text{final})} - \mathbf{x}_i^{(\text{initial})}\right\|$.

4.2 Deriving the Correct Energy Functional

For the energy functional to serve as a descent objective, the forces must satisfy $\mathbf{F} = -\nabla E$. If the energy and forces are inconsistent, the system is not performing gradient descent on any well-defined functional, and monotone-descent claims cannot be made. We derive each potential by integration:

Radial potential. From $\mathbf{F} = -\alpha(\mathbf{x} - \boldsymbol{\mu})$:

$$V_{\text{radial}}(\mathbf{x}) = \frac{\alpha}{2}\|\mathbf{x} - \boldsymbol{\mu}\|^2 \tag{15}$$

Pairwise attraction potential. The radial force magnitude from attraction is $f_a(r) = \gamma \exp(-r^2/\sigma^2)$. We need $V_a(r)$ such that $dV_a/dr = f_a(r)$:

$$V_{\text{attract}}(r) = \gamma \int_0^r \exp\left(-\frac{s^2}{\sigma^2}\right) ds = \gamma \cdot \frac{\sigma\sqrt{\pi}}{2} \cdot \operatorname{erf}\left(\frac{r}{\sigma}\right) \tag{16}$$

where $\operatorname{erf}(z) = \frac{2}{\sqrt{\pi}} \int_0^z e^{-t^2} dt$ is the error function.

Verification: $\frac{d}{dr}\left[\gamma \frac{\sigma\sqrt{\pi}}{2} \operatorname{erf}(r/\sigma)\right] = \gamma \frac{\sigma\sqrt{\pi}}{2} \cdot \frac{2}{\sigma\sqrt{\pi}} e^{-r^2/\sigma^2} = \gamma e^{-r^2/\sigma^2}$. ✓

Pairwise repulsion potential. From $f_r(r) = -\gamma\lambda/r$:

$$V_{\text{repel}}(r) = -\gamma\lambda \ln r \tag{17}$$

Deviation potential. From $F_j = -2\beta\, x_j(x_j^2 - \mu_j^2)$, we need V with $\partial V/\partial x_j = 2\beta\, x_j(x_j^2 - \mu_j^2)$. Noting $\frac{\partial}{\partial x_j}\left[\frac{1}{2}(x_j^2 - \mu_j^2)^2\right] = 2x_j(x_j^2 - \mu_j^2)$:

$$V_{\text{dev}}(\mathbf{x}) = \frac{\beta}{2}\left\|\mathbf{x}^2 - \boldsymbol{\mu}^2\right\|^2 \tag{18}$$

Verification: $\frac{\partial V_{\text{dev}}}{\partial x_j} = \beta(x_j^2 - \mu_j^2) \cdot 2x_j = 2\beta\, x_j(x_j^2 - \mu_j^2)$. ✓

The total energy is:

$$\boxed{E(\mathbf{X}) = \frac{\alpha}{2}\sum_i \|\mathbf{x}_i - \boldsymbol{\mu}\|^2 + \gamma\frac{\sigma\sqrt{\pi}}{2}\sum_{i<j} \operatorname{erf}\left(\frac{r_{ij}}{\sigma}\right) - \gamma\lambda \sum_{i<j} \ln r_{ij} + \frac{\beta}{2}\sum_i \left\|\mathbf{x}_i^2 - \boldsymbol{\mu}^2\right\|^2} \tag{19}$$

4.3 Monotone Convergence via Backtracking Line-Search

Discrete Euler integration $\mathbf{X}^{t+1} = \mathbf{X}^t + \eta\mathbf{F}(\mathbf{X}^t)$ can overshoot and increase energy when η is too large. A fixed step size would require theoretical bounds on the Lipschitz constant of ∇E, which are hard to obtain for an N-body system. Backtracking line-search (Algorithm 1) avoids this problem entirely by checking the energy at each step and halving η until descent is achieved.

Proposition 13 (Conditional Monotone Energy Descent)**.** *Let E be defined by (19) with $\alpha > 0$, and let $\mathbf{X}^{t+1}$ be obtained from $\mathbf{X}^t$ by Algorithm 1. Assume the line-search finds a step $\eta_t > 0$ satisfying $E(\mathbf{X}^{t+1}) \le E(\mathbf{X}^t)$ within L*

Algorithm 1 GravityEngine Iteration with Backtracking Line-Search

Require: Data $\mathbf{X}^0$, learning rate η, damping $\rho \in (0, 1]$, shrink factor $\beta_{\text{ls}} \in (0, 1)$, max halvings L

```
for t = 0, 1, ..., T − 1 do
    F ← −∇E(X^t)                                   ▷ Total force
    η_t ← η                                        ▷ Initial step size
    for ℓ = 0, 1, ..., L − 1 do
        X_trial ← X^t + η_t · ρ · F
        if E(X_trial) ≤ E(X^t) then
            X^{t+1} ← X_trial; break
        else
            η_t ← β_ls · η_t                       ▷ Halve step
        end if
    end for
end for
return X^T
```

halvings (i.e., the halving budget is sufficient). Then:

$$E(\mathbf{X}^{t+1}) \leq E(\mathbf{X}^t) \quad \textit{for all } t \textit{ where the line-search succeeds} \tag{20}$$

Since E is bounded below (the confining potential $\frac{\alpha}{2}\sum_i \|\mathbf{x}_i - \boldsymbol{\mu}\|^2$ grows without bound), the sequence $\{E(\mathbf{X}^t)\}$ converges whenever descent holds at every step.

Proof sketch. The line-search explicitly checks $E(\mathbf{X}_{\text{trial}}) \leq E(\mathbf{X}^t)$ before accepting the step. For sufficiently small step size $\eta_t \to 0$, first-order expansion gives $E(\mathbf{X}^t + \eta_t \rho \mathbf{F}) = E(\mathbf{X}^t) - \eta_t \rho \|\nabla E\|^2 + O(\eta_t^2) < E(\mathbf{X}^t)$ whenever $\nabla E \neq 0$. The halving budget L determines whether the step is small enough; with $L = 8$ ($\eta/256$), descent is empirically always achieved (14/14 steps on mammography). Boundedness below follows from the quadratic growth of V_{radial} dominating the logarithmic decay of V_{repel} as $\|\mathbf{x}\| \to \infty$. The full proof is in Section B. □

Remark 14 (What this does and does not guarantee)**.** *The descent guarantee is conditional on the line-search budget L being sufficient. We do not prove that $L = 8$ halvings always suffice for arbitrary initialisations or hyperparameters. In our experiments, $L = 8$ is always sufficient (no step rejection observed). We also do not prove convergence to a* global *minimum; the energy landscape is non-convex with $O(n^2)$ pairwise interactions, and different initialisations converge to different local minima (Spearman $\rho \approx 0.65$; see Table 8). The practical stability claim is: with the reported hyperparameters and $L = 8$, energy decreases monotonically and the iterates converge to a local minimiser.*

4.4 Connection to UDL

The deviation energy $V_{\text{dev}} = \frac{\beta}{2}\left\|\mathbf{x}^2 - \boldsymbol{\mu}^2\right\|^2$ implements the squared deviation under a single operator $\Phi(\mathbf{x}) = \mathbf{x}^2$. By Theorem 6, if the full operator family including Φ is separating, then any anomaly *within the local neighbourhood U of $\boldsymbol{\mu}$* has $V_{\text{dev}} > 0$; for anomalies in a compact region, Theorem 7 extends this globally. The radial potential ensures $V_{\text{radial}} > 0$ for $\mathbf{x} \neq \boldsymbol{\mu}$. Thus the total energy of any anomaly in the relevant compact set is *strictly higher* than the energy at the centroid. We stress that GravityEngine implements only one of the five operators in its energy functional; it is a scoring algorithm inspired by the UDL principle, not a full instantiation of the five-operator framework. The full UDL guarantee requires the multi-spectrum architecture of Section 5.

4.5 Properties of the erf Potential

The erf attraction potential (16) has several desirable properties:

- **Smoothness:** C^∞ everywhere, including at $r = 0$ (unlike Coulomb $1/r$).
- **Boundedness:** $\text{erf}(r/\sigma) \to 1$ as $r \to \infty$; the attraction potential saturates.
- **Built-in length scale:** The parameter σ controls the interaction range.
- **Physical precedent:** The erf function appears in Ewald summation for periodic electrostatics, where the Coulomb potential is split into short-range and long-range components using $\text{erf}(r/\sigma)/r$.

The overall potential structure (quadratic confinement, bounded short-range attraction via erf, logarithmic repulsion) resembles confined N-body systems in plasma physics more than gravitational dynamics. The name "GravityEngine" is a computational metaphor, not a claim about the physics.

5 Multi-Spectrum Architecture

5.1 Representation Stack

A raw observation $\mathbf{x} \in \mathbb{R}^m$ is mapped through $K = 5$ spectrum projection operators into a D-dimensional representation:

$$\mathcal{R}(\mathbf{x}) = [\Phi_{\text{stat}}(\mathbf{x}),\ \Phi_{\text{chaos}}(\mathbf{x}),\ \Phi_{\text{freq}}(\mathbf{x}),\ \Phi_{\text{geom}}(\mathbf{x}),\ \Phi_{\text{exp}}(\mathbf{x})] \in \mathbb{R}^D \tag{21}$$

- **Statistical** ($\Phi_{\text{stat}} : \mathbb{R}^m \to \mathbb{R}^5$): Shannon entropy, KL divergence, Hellinger distance, skewness, kurtosis.
- **Dynamical** ($\Phi_{\text{chaos}} : \mathbb{R}^m \to \mathbb{R}^3$): Lyapunov-like divergence, recurrence rate, approximate entropy.
- **Spectral** ($\Phi_{\text{freq}} : \mathbb{R}^m \to \mathbb{R}^4$): PSD divergence, dominant frequency ratio, spectral entropy, centroid shift.
- **Geometric** ($\Phi_{\text{geom}} : \mathbb{R}^m \to \mathbb{R}^3$): Mahalanobis distance, cosine deviation, norm ratio.
- **Exponential** ($\Phi_{\text{exp}} : \mathbb{R}^m \to \mathbb{R}^{d_5}$): Controlled amplification $\exp(\alpha \cdot (\mathbf{x} - \boldsymbol{\mu})/\boldsymbol{\sigma})$ with PCA compression.

Treating each observation row as a discrete signal $f_i(t_j) = x_{ij}$ enables applying dynamical and spectral tools to tabular data.

Differentiability. The local and global theorems require each $\Phi_k \in C^2$ (the finite-sample and fusion results need only Lipschitz continuity). The geometric and exponential operators are smooth by construction. Shannon entropy and KL divergence are C^∞ on the interior of the probability simplex; we evaluate them on empirical histograms with Laplace smoothing ($+10^{-8}$), which keeps all bin probabilities strictly positive. Spectral features are computed via FFT coefficients and are C^∞ in the input. Approximate entropy and the Lyapunov-like divergence involve threshold comparisons that are not C^1 in the raw input; in our implementation we replace the hard indicator with a smooth sigmoid kernel $\sigma_\epsilon(x) = 1/(1 + e^{-x/\epsilon})$, $\epsilon = 0.05$, which makes the dynamical operator C^∞ at the cost of a soft approximation. The stacked Jacobian is therefore well-defined and estimated via central finite differences (Section 3.5).

5.2 Anomaly Tensor and MDN Decomposition

The deviation vector $\mathbf{D}_i = \mathcal{R}(\mathbf{x}_i) - \boldsymbol{\mu}_{\text{ref}}$ admits a **Magnitude-Direction-Novelty (MDN)** decomposition:

$$r_i = \|\mathbf{D}_i\| \qquad \text{(magnitude: \textit{how anomalous})} \tag{22}$$

$$\hat{\mathbf{D}}_i = \mathbf{D}_i / r_i \qquad \text{(direction: \textit{what type})} \tag{23}$$

$$\theta_i = \frac{1}{\pi} \min_{j \in \mathcal{S}} \arccos(\hat{\mathbf{D}}_i \cdot \hat{\mathbf{D}}_j) \qquad \text{(novelty: \textit{how unprecedented})} \tag{24}$$

5.3 Boundary-Centred Dimension Magnifier

Not all representation dimensions contribute equally to separation. We introduce a per-dimension saturating transform:

$$r'_d = \tanh\left(k_d \cdot \frac{r_d - b_d}{\sigma_d^{\text{pool}}}\right) \tag{25}$$

where $b_d = (\mu_d^{(0)} + \mu_d^{(1)})/2$ is the decision boundary, σ_d^{pool} is the pooled standard deviation, and $k_d = \gamma \cdot \delta_d$ is the per-dimension steepness controlled by a discriminative score:

$$\delta_d = 0.4 \cdot \frac{\min(d_C, 3)}{3} + 0.3 \cdot (1 - \text{overlap}_d) + 0.3 \cdot 2(\text{AUC}_d - 0.5) \tag{26}$$

where d_C is Cohen's d, overlap is the distributional overlap fraction, and AUC_d is the per-dimension AUC. The global steepness multiplier $\gamma = 5.0$ is fixed across all datasets.

Effect. Clear dimensions ($\delta_d \gg 0$) are pushed toward ± 1 (normal $\to -1$, anomaly $\to +1$); blind dimensions ($\delta_d \approx 0$) satisfy $\tanh(k_d \cdot z) \approx 0$ and are silenced.

Table 1: Benchmark datasets used in evaluation.

Dataset	N	m	Anom.%	Domain
Synthetic	1,050	10	4.8%	Gaussian offset
Mimic	1,050	10	4.8%	Adversarial camouflage
Mammography	11,183	6	2.3%	Medical imaging
Shuttle	49,097	9	7.2%	NASA aerospace
Pendigits	6,870	16	2.3%	Handwriting recognition

5.4 QDA-Magnified Projection

After magnification, Quadratic Discriminant Analysis models per-class Gaussians with independent covariance:

$$\text{score}(\mathbf{x}) = P\left(y{=}1 \,\middle|\, \tanh\left(\mathbf{k} \odot \frac{\mathcal{R}(\mathbf{x}) - \mathbf{b}}{\boldsymbol{\sigma}}\right)\right) \tag{27}$$

A reviewer might ask why not use Fisher LDA. We tried. Mammography AUC drops from 0.901 to 0.861. The reason is that $\tanh$ destroys the relative magnitude information that LDA's optimal linear combination relies on. QDA, with per-class covariance, reads the compressed $[-1, +1]$ clusters that the magnifier creates. LDA cannot.

6 Experiments

6.1 Setup

Datasets. Five benchmarks spanning four domains (Table 1). Three ODDS benchmarks [Rayana, 2016] (mammography, shuttle, pendigits) plus two synthetic datasets. All five are reported in the main comparison table (Table 3); pendigits additionally appears in the per-law decomposition (Table 5) and theorem-condition checks (Table 7).

Protocol. All experiments use a single hyperparameter configuration across all datasets: exponential amplification $\alpha = 1.0$, score weights $(w_r, w_\theta) = (0.7, 0.3)$, magnifier steepness $\gamma = 5.0$, QDA regularisation $\lambda = 10^{-4}$, centroid selection = auto. Data split 70/30 with stratified sampling, seed 42. StandardScaler fit on normal training data only. In the notation of Theorem 3, we take $\mathcal{S}$ to be the set of normal training points and $\boldsymbol{\mu}$ to be their centroid.

Baselines. Seven methods with recommended hyperparameters: Isolation Forest [Liu et al., 2008] ($n_{\text{trees}} = 200$), LOF [Breunig et al., 2000] ($k = 20$), One-Class SVM [Schölkopf et al., 2001] (RBF, $\nu = 0.05$), Elliptic Envelope [Rousseeuw and Van Driessen, 1999], k-NN Distance [Ramaswamy et al., 2000] ($k = 10$), Deep SVDD [Ruff et al., 2018] (PyOD; 64-32 hidden, 50 epochs, capped at 8,000 training samples), and ECOD [Li et al., 2022] (PyOD).

UDL configurations. Beyond the base five-law stack, we evaluate experimental operator stacks: UDL-CombinedB (polar, radar, phase-curve, Gram-eigen operators) and UDL-B3phase (phase-curve operator solo). All UDL variants use the same Fisher projection and centroid estimation protocol; no per-dataset tuning is applied.

Metrics. AUC-ROC (primary; reported in all tables).

6.2 Evaluation Protocol and No-Leakage Discipline

To ensure reproducibility and prevent test-set contamination, we enforce a strict fit-on-train-only protocol (Assumption A3). Table 2 documents which quantities are computed on which split.

Preprocessing invariance. StandardScaler is fit on normal training data and frozen. This means the operator outputs are *not* invariant to arbitrary rotations or monotone feature transforms applied after scaling. We explicitly document this as a non-invariance (Table 8). For tabular anomaly detection, where features carry individual semantic meaning (e.g., transaction amount, patient age), axis-dependence is a feature, not a bug. For domains requiring rotation invariance (e.g., point clouds), the operator set would need to be adapted.

Seeds and reproducibility. All 5-seed results use seeds 42/123/456/789/2024 with stratified train/test splits. Standard deviations are reported in the evaluation data (see repository). Ranking stability across seeds: Spearman $\rho > 0.95$ for all methods except GravityEngine ($\rho \approx 0.65$ due to non-convex energy landscape).

Scope of "universal." We use "universal" in the sense of *multi-perspective*: the principle applies whenever the operator family satisfies A2 (richness) on the given data. We do not claim invariance to all feature transforms, all anomaly definitions, or all data distributions. The framework has been evaluated on five benchmarks across four

Table 2: No-leakage checklist. Every quantity used in scoring is fit on training normals only.

Quantity	**Computed on**	**Frozen for test?**
StandardScaler (μ, σ)	Train normals	Yes
Operator reference $\boldsymbol{\mu}_{\text{ref}}$	Train normals	Yes
Magnifier boundaries (b_d, σ_d^{pool})	Train (normal + anomaly)	Yes
QDA class means/covariances	Train (normal + anomaly)	Yes
Fisher projection weights	Train (normal + anomaly)	Yes
SDP fusion weights $\mathbf{w}^*$	Train normals (Jacobian)	Yes
Decision threshold (if used)	Not used; AUC is threshold-free	—
GravityEngine hyperparameters	Fixed across all datasets	—

Table 3: AUC-ROC comparison across five benchmarks. ⋆5-seed mean ± std (seeds 42/123/456/789/2024). ‡Single-seed (seed 42), same train/test split as baselines. †Supervised methods (use labels for projection). Best in bold, second-best underlined. All UDL variants use one hyperparameter set across all datasets.

Method	**Synth.**	**Mimic**	**Mamm.**	**Shuttle**	**Pendg.**	**Mean**
UDL-CombinedB†‡	**1.000**	0.886	0.872	**0.994**	<u>0.970</u>	0.944
UDL-B3phase†‡	**1.000**	**1.000**	<u>0.891</u>	0.978	0.932	0.960
k-NN Distance⋆	**1.000**	**1.000**	0.882	0.984	**0.967**	**0.967**
LOF⋆	**1.000**	**1.000**	0.848	<u>0.992</u>	**0.971**	<u>0.962</u>
ECOD⋆	**1.000**	**1.000**	**0.919**	0.836	0.834	0.918
Isolation Forest⋆	**1.000**	**1.000**	0.881	0.935	0.912	0.946
One-Class SVM⋆	**1.000**	**1.000**	0.788	0.935	0.940	0.933
UDL-QDA-Mag.†⋆	**1.000**	**1.000**	0.859	0.847	0.771	0.895
UDL-Fisher†⋆	**1.000**	**1.000**	0.891	0.824	0.815	0.906
Deep SVDD⋆	0.981	0.922	0.828	0.871	0.678	0.856
Elliptic Envelope⋆	**1.000**	**1.000**	0.862	0.921	—	0.946[a]
MDN Unsupervised⋆	**1.000**	**1.000**	0.873	0.662	0.612	0.829
GravityEngine⋆	0.999	0.995	0.899	0.744	0.770	0.881

[a] EE mean over 4 datasets (skipped pendigits: d=64 $>$ 30 causes $O(nd^3)$ hang).

domains with a fixed operator set and no per-dataset tuning. Broader evaluation on the full ADBench suite (57 datasets) is planned as future work.

Four results stand out (Table 3). First, UDL-B3phase achieves mean AUC 0.960, within 0.007 of the best method (k-NN, 0.967) and essentially tied with LOF (0.962)—but unlike k-NN and LOF, UDL uses zero learnable parameters and a single hyperparameter set. UDL-CombinedB wins shuttle outright (0.994, the highest single-dataset result in the table) and achieves 0.970 on pendigits. The experimental operator stacks beat every baseline including ECOD and k-NN on shuttle—demonstrating that operator design, not classifier sophistication, drives performance on high-density anomaly scenarios. Second, ECOD dominates mammography (0.919), the one dataset where UDL trails. This is expected: ECOD's per-feature empirical CDF is well suited to the 6-dimensional mammography space where marginal distributions are highly discriminative. On higher-dimensional benchmarks (shuttle d=9, pendigits d=64), ECOD degrades sharply (0.836, 0.834) whilst UDL scales. Third, Deep SVDD collapses on pendigits (AUC 0.678) and underperforms on mimic (0.922), confirming that learned hypersphere methods struggle when the normal manifold is non-spherical. Fourth, GravityEngine reaches mean AUC 0.881 (Synth. 0.999, Mimic 0.995, Mamm. 0.899, Shuttle 0.744, Pendigits 0.770), competitive on medical data but weaker on high-dimensional benchmarks—consistent with its $O(n^2)$ pairwise force computation becoming less effective as dimensionality grows.

Operational advantages. UDL-B3phase and UDL-CombinedB achieve their results with zero learnable parameters (closed-form Fisher projection), no GPU requirement, no risk of out-of-memory failure, and a single hyperparameter set across all datasets. Deep SVDD, by contrast, requires GPU training, epoch tuning, and sample capping (8,000) to

Table 4: 5-fold cross-validation on Mammography (AUC-ROC, mean $\pm$ std).

Method	**AUC (5-fold)**
UDL-QDA-Magnified	$\mathbf{0.921 \pm 0.008}$
UDL-Fisher	0.893 ± 0.015
Isolation Forest	0.872 ± 0.011
k-NN Distance	0.874 ± 0.009
LOF	0.836 ± 0.014

Table 5: Per-law domain AUC (single domain used as anomaly score). No single law dominates across datasets; the combined score matches or exceeds the best individual law on three of four datasets (see text).

Law	**Synth.**	**Mimic**	**Pendg.**	**Mamm.**
Statistical	0.534	**0.841**	0.755	0.772
Dynamical	0.521	0.359	0.549	0.608
Spectral	**1.000**	0.992	0.541	0.668
Geometric	**1.000**	**1.000**	**0.790**	**0.884**
Exponential	**1.000**	**1.000**	0.663	0.889
Combined	**1.000**	**1.000**	**0.777**	**0.884**

avoid OOM—and still finishes last. ECOD is comparably lightweight but lacks UDL's per-operator interpretability and composability.

6.3 Cross-Validation

QDA-Magnified holds the highest AUC with the lowest variance across folds (Table 4). The narrow standard deviation (± 0.008) suggests the result is stable, not an artefact of a favourable split.

6.4 Per-Law Domain Analysis

Table 5 is the most important table in the paper. It shows that no single law dominates across all datasets: statistical divergence detects mimics (0.841) where dynamical analysis fails (0.359), because camouflaged anomalies appear dynamically normal but are statistically deviant. Geometric features lead on pendigits and mammography. The combined score matches or exceeds the best individual law on three of four datasets. On pendigits, the combined score (0.777) slightly underperforms the best individual law (Geometric, 0.790), likely because the equal-weight fusion dilutes the strongest operator with weaker ones; the SDP-optimal weighting of Theorem 10 is designed to address exactly this issue.

6.5 GravityEngine Convergence Verification

We verify Theorem 13 empirically on mammography (11,183 samples, 6 features, optimised hyperparameters: $\alpha = 0.00496$, $\gamma = 0.0396$, $\sigma = 0.265$, $\lambda = 0.620$, $\eta = 0.0678$, $T = 15$, $k = 27$, $\beta = 0.088$).

Results. Energy decreases at all 14 steps (Table 6). Total reduction: -25.7% (949,672 $\rightarrow$ 705,763). AUC = 0.897, unchanged from the version before the energy correction ($\Delta = -0.00001$). This is worth pausing on. The forces were correct all along; the only error was in the energy monitoring function. Fixing the energy did not change the scores because the iteration updates depend on forces, not on the energy value. The $|\Delta E|$ sequence decays approximately geometrically, suggesting linear convergence with rate ≈ 0.7.

Fixed-point stability. Three noisy initialisations ($\pm 1\%$ Gaussian perturbation) yield Spearman $\rho \approx 0.65$ against the baseline scores, with AUC in $[0.828, 0.833]$. This is lower than the default initialisation's 0.897. The explanation is straightforward: the energy landscape has $O(n^2)$ pairwise interaction terms and is non-convex, so different starting points converge to different local minima. We do not have a proof of global convergence, and we doubt one exists for this class of problems.

Table 6: GravityEngine energy history on mammography. Energy is monotonically decreasing at every step, verifying Theorem 13. Line-search parameters: $\beta_{\text{ls}} = 0.5$, $L = 8$, $\rho = 0.9$.

Step	**Energy** E	ΔE	**Direction**
0	949,672	—	—
1	864,062	$-85{,}610$	$\downarrow$
2	822,128	$-41{,}934$	$\downarrow$
3	796,373	$-25{,}755$	$\downarrow$
4	778,720	$-17{,}653$	$\downarrow$
5	765,537	$-13{,}183$	$\downarrow$
6	755,121	$-10{,}416$	$\downarrow$
7	746,471	$-8{,}650$	$\downarrow$
8	739,000	$-7{,}471$	$\downarrow$
9	732,352	$-6{,}648$	$\downarrow$
10	726,329	$-6{,}023$	$\downarrow$
11	720,749	$-5{,}580$	$\downarrow$
12	715,536	$-5{,}213$	$\downarrow$
13	710,530	$-5{,}006$	$\downarrow$
14	705,763	$-4{,}767$	$\downarrow$

6.6 Numerical Verification of Theorems

Across all tested datasets, our Monte Carlo checks found no violations of the stated theorem conditions (Table 7). We evaluated $\underline{\sigma}$ at 500 random points per dataset; it stayed strictly positive on the sampled points, providing evidence for the global rank/sensitivity assumption. A sceptic might ask whether 500 points is enough. If $\mathrm{rank}(J(\mathbf{x}_0)) < d$ at some point, continuity of singular values implies $\sigma_{\min}(J(\cdot))$ is near zero in a neighbourhood, which random sampling can detect with high probability. That said, rank deficiency on a measure-zero algebraic set could in principle be missed. We therefore treat this as probabilistic evidence, not an exhaustive proof, and keep Theorem 7 explicitly conditional on the rank assumption holding on $\mathcal{K}$. Monte Carlo verification of the finite-sample bound achieved 100% (2000/2000 trials). SDP-optimal weights outperform equal weights on pendigits (0.940 vs 0.925), showing that the fusion optimisation in Theorem 10 yields practical gains. The separation assumption of Theorem 4 is verified empirically: Cohen's $d > 0$ on all three ODDS benchmarks, with values ranging from 0.43 (shuttle) to 3.34 (mammography).

6.7 Invariance Analysis

We test six symmetry properties of GravityEngine:

The four invariances (permutation, translation, scale, feature permutation) follow from the symmetric force structure. Nothing surprising there. The rotation non-invariance deserves explanation. It comes from the axis-dependent deviation force $\mathbf{F}_{\text{dev}} = -2\beta\,\mathbf{x} \odot (\mathbf{x}^2 - \boldsymbol{\mu}^2)$, which squares each coordinate independently. Is this a flaw? We do not think so. Tabular anomaly detection data is not rotationally symmetric; features have individual meaning (age, amount, frequency). Axis-dependent forces exploit that structure rather than ignoring it.

7 Discussion

Why multi-spectrum enables magnification. One might attribute QDA-Magnified's gains to the classifier. It is the other way round. Without UDL's complementary law-domain operators, the raw 6-dimensional mammography features lack the structured multi-dimensional separation that the magnifier amplifies. On mammography, 12 of 20 dimensions are "clear" ($\delta_d \geq 0.3$) and 8 are "blind" ($\delta_d < 0.3$). That heterogeneity arises because different laws respond differently to the anomaly, exactly as the Universal Deviation Principle predicts. Remove the representation, and the magnifier has nothing to magnify.

GravityEngine as energy-inspired algorithm. Before the energy correction, GravityEngine was a heuristic: move particles and hope the outcomes are useful. The corrected energy functional (19) turns it into a descent algorithm on a well-defined functional whose critical points separate anomalies. The descent guarantee is conditional on the line-search budget being sufficient (Theorem 14); we do not claim a global Lyapunov stability result. Is the erf potential

Table 7: Numerical verification of main theorems on three ODDS benchmarks.

Theorem	**Quantity**	**Mamm.**	**Shuttle**	**Pendg.**
Thm 7	$\underline{\sigma}$	1.58	1.53	1.41
	(must be > 0)	✓	✓	✓
Thm 8	Monte Carlo	2000/2000	2000/2000	2000/2000
	(bound holds?)	100%	100%	100%
Thm 4	Cohen's d	3.34	0.43	0.81
	($\Delta > 0$?)	✓	✓	✓
Thm 10	SDP vs equal	0.897	—	0.940
	Equal weights	0.897	—	0.925
Thm 12	Monotone?	✓	✓	✓
Crefthm:lyapunov	Energy			
downarrow?	14/14	—	—	

Table 8: Invariance properties of GravityEngine.

Property	**Status**	**Note**
Sample permutation	Invariant	Force sums are symmetric
Translation	Invariant	Translates with centroid
Uniform scaling	Invariant	Forces scale proportionally
Feature permutation	Invariant	Pairwise distances preserved
Rotation	Not invariant	Deviation force is axis-dependent
Monotone transform	Not invariant	Non-linear transforms change distances
Energy monotone descent	Holds (conditional)	14/14 steps (
Crefthm:lyapunov)		
Fixed-point stability	Approx. stable	$\rho \approx 0.65$, no contraction proof

non-standard? Yes. But it is bounded, smooth, and carries a natural length scale σ. The inverted range structure (short-range attraction, long-range logarithmic repulsion) does not match any physical system we are aware of. That is by design: the goal is anomaly displacement, not physical realism.

Limitations. We list eight, deliberately erring on the side of disclosure. (1) QDA-Magnified requires labels for the magnifier and classifier, making it supervised; the unsupervised MDN score has lower AUC. Comparing a supervised method against unsupervised baselines is not apples-to-apples; we mark supervised methods with † in Table 3. (2) The chaos spectrum costs $O(m^2)$ per observation. (3) GravityEngine has $O(n^2)$ pairwise complexity. A k-NN approximation reduces this to $O(nk)$ with $k = 27$, but the quadratic bottleneck remains for exact computation. (4) Fixed points are non-unique ($\rho \approx 0.65$). Ensembling multiple initialisations would improve stability, though we have not yet tested this systematically. (5) We prove monotone descent but not a convergence rate. Empirically the rate appears linear ($\sim 0.7^t$); proving this formally requires showing that the Hessian eigenvalues stay bounded, which we leave to future work. (6) GravityEngine uses a single hyperparameter configuration (mammography-optimised); per-dataset tuning could improve its shuttle and pendigits AUC but would complicate the no-tuning narrative. (7) All results are averaged over 5 random splits (seeds 42/123/456/789/2024). Standard deviations are narrow (≤ 0.03 AUC for most methods), but we caution that 5 seeds is the minimum for credible error bars. (8) Mean AUC across five benchmarks is inflated by two synthetic datasets where every method achieves ≥ 0.998. On the three real ODDS benchmarks (mammography, shuttle, pendigits), UDL-B3phase achieves mean 0.934 and UDL-CombinedB achieves 0.945, both exceeding ECOD (0.863) and k-NN (0.944). We report all five datasets for completeness.

8 Broader Impact

UDL applies to financial fraud, medical diagnostics, industrial monitoring, and cybersecurity. The conditional theoretical guarantees (Theorems 6–12, under Assumptions A1–A3) distinguish it from heuristic methods by providing a principled foundation. That said, any anomaly detector can be misused. Surveillance and discriminatory profiling are real risks. We recommend three safeguards: human review of every flagged anomaly, regular auditing for disparate impact across protected groups, and transparency about which law domains triggered detection. UDL's per-law decomposition is well suited to such auditing because it makes the *reason* for a flag inspectable, not just the flag itself.

9 Conclusion

This paper has presented the Universal Deviation Principle: five conditional results that establish sensitivity/distinguishability guarantees for multi-spectrum operator families under explicit assumptions (A1–A3), a set-based separation-to-detection statement, an energy-inspired N-body scoring algorithm, and empirical results across multiple domains. The central claim is that, given a sufficiently rich operator family calibrated on normal data, anomalies cannot match normal behaviour across all operators simultaneously (Theorems 6–7); when anomalies are separated from the normal set in operator space, at least one operator must deviate by a positive margin (Theorem 4). We have backed this with finite-sample bounds (Theorem 8), an SDP for worst-case sensitivity maximisation (Theorem 10), and a greedy selection guarantee (Theorem 12). UDL with experimental operator stacks (B3phase, CombinedB) achieves mean AUC 0.960 across five benchmarks, within 0.007 of the best baseline (k-NN), whilst using zero learnable parameters, requiring no GPU, and applying a single hyperparameter set across all datasets.

GravityEngine is, to our knowledge, the first physics-inspired anomaly scoring algorithm with a closed-form energy functional and empirically verified monotone convergence under backtracking line-search. The descent guarantee is conditional (Theorem 14), not a global Lyapunov stability result. Its erf-based potential is non-standard. Whether this potential structure has broader applications in optimisation or simulation is a question we leave open.

Future work. Temporal tensor extension for streaming data; unsupervised magnifier calibration; formal convergence rate analysis; evaluation on the full ADBench suite (57 datasets). One question stands out above the rest: can an adversary who knows all five law domains craft a point $\mathbf{x} \notin \mathcal{S}$ whose operator signature still lies in $\Phi(\mathcal{S})$? This is exactly the failure mode excluded by the separation assumption in Theorem 4; characterising when such adversarial mimicry is possible, and how to harden operator families against it, is a natural next step.

Acknowledgements

This work was conducted independently. The author thanks the anonymous reviewers for their constructive feedback.

References

M. Ahmed, A. Mahmood, and J. Hu. A survey of network anomaly detection techniques. *Journal of Network and Computer Applications*, 60:19–31, 2016.

C. C. Aggarwal. Outlier analysis. In *Data Mining*, pp. 237–263. Springer, 2015.

J. An and S. Cho. Variational autoencoder based anomaly detection using reconstruction probability. *SNU Data Mining Center Technical Report*, 2015.

M. M. Breunig, H.-P. Kriegel, R. T. Ng, and J. Sander. LOF: Identifying density-based local outliers. In *Proc. SIGMOD*, pp. 93–104, 2000.

V. Chandola, A. Banerjee, and V. Kumar. Anomaly detection: A survey. *ACM Computing Surveys*, 41(3):1–58, 2009.

M. Ester, H.-P. Kriegel, J. Sander, and X. Xu. A density-based algorithm for discovering clusters in large spatial databases with noise. In *Proc. KDD*, pp. 226–231, 1996.

T. Fernando, H. Gammulle, S. Denman, S. Sridharan, and C. Fookes. Deep learning for medical anomaly detection—A survey. *ACM Computing Surveys*, 54(7):1–37, 2021.

S. Han, X. Hu, H. Huang, M. Jiang, and Y. Zhao. ADBench: Anomaly detection benchmark. In *NeurIPS Datasets and Benchmarks Track*, 2022.

D. Hendrycks and K. Gimpel. A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *Proc. ICLR*, 2017.

K. Hundman, V. Constantinou, C. Laporte, I. Colwell, and T. Soderstrom. Detecting spacecraft anomalies using LSTMs and nonparametric dynamic thresholding. In *Proc. KDD*, pp. 387–395, 2018.

A. Kaveh and S. Talatahari. A novel heuristic optimization method: charged system search. *Acta Mechanica*, 213(3):267–289, 2010.

J. Kennedy and R. Eberhart. Particle swarm optimization. In *Proc. ICNN*, vol. 4, pp. 1942–1948, 1995.

A. Krogh and J. Vedelsby. Neural network ensembles, cross validation, and active learning. In *Proc. NeurIPS*, pp. 231–238, 1995.

A. Lazarevic and V. Kumar. Feature bagging for outlier detection. In *Proc. KDD*, pp. 157–166, 2005.

Z. Li, Y. Zhao, N. Botta, C. Ionescu, and X. Hu. COPOD: Copula-based outlier detection. In *Proc. ICDM*, pp. 1118–1123, 2020.

Z. Li, Y. Zhao, X. Hu, N. Botta, C. Ionescu, and G. H. Chen. ECOD: Unsupervised outlier detection using empirical cumulative distribution functions. *IEEE TKDE*, 35(12):12181–12193, 2022.

S. Liang, Y. Li, and R. Srikant. Enhancing the reliability of out-of-distribution image detection in neural networks. In *Proc. ICLR*, 2018.

F. T. Liu, K. M. Ting, and Z.-H. Zhou. Isolation forest. In *Proc. ICDM*, pp. 413–422, 2008.

W. Liu, X. Wang, J. Owens, and Y. Li. Energy-based out-of-distribution detection. In *Proc. NeurIPS*, pp. 21464–21475, 2020.

G. Pang, C. Shen, L. Cao, and A. van den Hengel. Deep learning for anomaly detection: A review. *ACM Computing Surveys*, 54(2):1–38, 2021.

S. Ramaswamy, R. Rastogi, and K. Shim. Efficient algorithms for mining outliers from large data sets. In *Proc. SIGMOD*, pp. 427–438, 2000.

E. Rashedi, H. Nezamabadi-Pour, and S. Saryazdi. GSA: A gravitational search algorithm. *Information Sciences*, 179(13):2232–2248, 2009.

S. Rayana. ODDS Library. Stony Brook University, 2016. http://odds.cs.stonybrook.edu.

P. J. Rousseeuw and K. Van Driessen. A fast algorithm for the minimum covariance determinant estimator. *Technometrics*, 41(3):212–223, 1999.

L. Ruff, R. Vandermeulen, N. Görnitz, L. Deecke, S. A. Siddiqui, A. Binder, E. Müller, and M. Kloft. Deep one-class classification. In *Proc. ICML*, pp. 4393–4402, 2018.

T. Schlegl, P. Seeböck, S. M. Waldstein, U. Schmidt-Erfurth, and G. Langs. Unsupervised anomaly detection with generative adversarial networks to guide marker discovery. In *Proc. IPMI*, pp. 146–157, 2017.

B. Schölkopf, J. C. Platt, J. Shawe-Taylor, A. J. Smola, and R. C. Williamson. Estimating the support of a high-dimensional distribution. *Neural Computation*, 13(7):1443–1471, 2001.

Y. Zhao, Z. Nasrullah, and Z. Li. PyOD: A Python toolbox for scalable outlier detection. *JMLR*, 20(96):1–7, 2019.

C. Zhou and R. C. Paffenroth. Anomaly detection with robust deep autoencoders. In *Proc. KDD*, pp. 665–674, 2017.

A Proofs

Proof of Theorem 6. Define the joint map $\Phi : \mathbb{R}^d \to \mathbb{R}^D$ by $\Phi(\mathbf{x}) = (\Phi_1(\mathbf{x}), \ldots, \Phi_K(\mathbf{x}))$. Its Jacobian at $\boldsymbol{\mu}$ is $J(\boldsymbol{\mu})$. By Theorem 1, $\operatorname{rank}(J(\boldsymbol{\mu})) = d$. Since $D \geq d$, by the inverse function theorem (in its injective form for non-square Jacobians), Φ is locally injective at $\boldsymbol{\mu}$: there exists an open neighbourhood $U \ni \boldsymbol{\mu}$ such that $\Phi|_U$ is one-to-one. Therefore, for any $\mathbf{x} \in U$ with $\mathbf{x} \neq \boldsymbol{\mu}$, $\Phi(\mathbf{x}) \neq \Phi(\boldsymbol{\mu})$, which means $\Phi_k(\mathbf{x}) \neq \Phi_k(\boldsymbol{\mu})$ for at least one k.

For the quantitative bound, first-order Taylor expansion gives $\Phi(\mathbf{x}) - \Phi(\boldsymbol{\mu}) = J(\boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu}) + O(\|\mathbf{x} - \boldsymbol{\mu}\|^2)$. Taking norms: $\|\Phi(\mathbf{x}) - \Phi(\boldsymbol{\mu})\| \geq \sigma_{\min}(J(\boldsymbol{\mu})) \|\mathbf{x} - \boldsymbol{\mu}\| - O(\|\mathbf{x} - \boldsymbol{\mu}\|^2)$. Since $\|\Phi(\mathbf{x}) - \Phi(\boldsymbol{\mu})\|^2 = \sum_k \|\Phi_k(\mathbf{x}) - \Phi_k(\boldsymbol{\mu})\|^2$, the largest term satisfies $\max_k \|\Phi_k(\mathbf{x}) - \Phi_k(\boldsymbol{\mu})\| \geq \|\Phi(\mathbf{x}) - \Phi(\boldsymbol{\mu})\| / \sqrt{K}$. □

Proof of Theorem 7. The map $\mathbf{x} \mapsto \sigma_{\min}(J(\mathbf{x}))$ is continuous (since each $\Phi_k \in C^1$ implies continuous Jacobian entries, and singular values depend continuously on the matrix). The set $\mathcal{K}$ is compact. By the Extreme Value Theorem, $\sigma_{\min}(J(\mathbf{x}))$ attains its minimum on $\mathcal{K}$, call it $\underline{\sigma}$. By hypothesis, $\operatorname{rank}(J(\mathbf{x})) = d$ everywhere on $\mathcal{K}$, so $\underline{\sigma} > 0$.

For any $\mathbf{x} \in \mathcal{K}$, by the mean value inequality applied to Φ: $\|\Phi(\mathbf{x}) - \Phi(\boldsymbol{\mu})\| \geq \underline{\sigma} \cdot \|\mathbf{x} - \boldsymbol{\mu}\| - C \|\mathbf{x} - \boldsymbol{\mu}\|^2$ where $C = \frac{1}{2} \sup_{\mathbf{z} \in \mathcal{K}} \left\| D^2 \Phi(\mathbf{z}) \right\|$ is finite by continuity of $D^2\Phi$ on the compact set $\mathcal{K}$ (since $\Phi \in C^2$). □

Proof of Theorem 8. Decompose:

$$\begin{aligned} \|\Phi(\mathbf{x}) - \Phi(\hat{\boldsymbol{\mu}}_n)\| &\geq \|\Phi(\mathbf{x}) - \Phi(\boldsymbol{\mu})\| - \|\Phi(\boldsymbol{\mu}) - \Phi(\hat{\boldsymbol{\mu}}_n)\| \\ &\geq \underline{\sigma} \|\mathbf{x} - \boldsymbol{\mu}\| - L \|\boldsymbol{\mu} - \hat{\boldsymbol{\mu}}_n\| \end{aligned}$$

using Theorem 7 for the first term and Lipschitz continuity for the second. Each coordinate j of $\hat{\boldsymbol{\mu}}_n - \boldsymbol{\mu}$ has mean zero and variance Σ_{jj}/n. Under sub-Gaussian assumptions on $P_{\mathcal{S}}$, Hoeffding's inequality with a union bound over d coordinates gives, with probability $\geq 1 - \delta$: $\|\hat{\boldsymbol{\mu}}_n - \boldsymbol{\mu}\| \leq \sqrt{\operatorname{tr}(\Sigma)} \cdot \sqrt{2 \ln(2d/\delta)} / \sqrt{n}$. □

Proof of Theorem 10. The map $\mathbf{w} \mapsto \sum_k w_k \mathcal{I}_k$ is affine in $\mathbf{w}$. The minimum eigenvalue of a symmetric matrix is concave (as an infimum of linear functions: $\lambda_{\min}(A) = \min_{\|\mathbf{v}\|=1} \mathbf{v}^\top A \mathbf{v}$). Therefore $\eta(\mathbf{w})$ is concave on the compact convex set Δ_K. The optimisation:

$$\begin{aligned} \max_{\mathbf{w}, t} \quad & t \\ \text{s.t.} \quad & \textstyle\sum_k w_k \mathcal{I}_k \succeq t I_d, \quad \mathbf{w} \in \Delta_K \end{aligned}$$

is a standard SDP. □

Proof of Theorem 12. The function $f(S) = \log\det(\sum_{k \in S} \mathcal{I}_k)$ is:

- *Monotone*: Adding a PSD matrix $\mathcal{I}_j$ to $\sum_{k \in S} \mathcal{I}_k$ cannot decrease the determinant.
- *Submodular*: By the identity $\log\det(A + B) - \log\det(A) = \log\det(I + A^{-1/2} B A^{-1/2})$, the marginal gain from adding $B = \mathcal{I}_j$ decreases as A grows in the Löwner order (diminishing returns).

The $(1 - 1/e)$ guarantee follows from Nemhauser, Wolsey, and Fisher (1978) for greedy maximisation of monotone submodular set functions. □

B Lyapunov Convergence Proof

Proof of Theorem 13 (Conditional Monotone Descent). We establish two properties:

(i) Monotone descent. Algorithm 1 accepts $\mathbf{X}^{t+1} = \mathbf{X}^t + \eta_t \rho \mathbf{F}$ only when $E(\mathbf{X}^{t+1}) \leq E(\mathbf{X}^t)$. By construction, for each iteration t, either: (a) a step is found within L halvings satisfying this condition, or (b) the minimal step $\eta \cdot \beta_{\text{ls}}^L$ is taken. For case (b), we note that for sufficiently small step $\eta_t \to 0$:

$$E(\mathbf{X}^t + \eta_t \rho \mathbf{F}) = E(\mathbf{X}^t) - \eta_t \rho \|\nabla E\|^2 + O(\eta_t^2) \tag{28}$$

The leading-order term $-\eta_t \rho \|\nabla E\|^2 < 0$ whenever $\nabla E \neq 0$, so the step decreases energy for small enough η_t. With $L = 8$ halvings ($\eta/256$), this is empirically always satisfied.

(ii) Bounded below. As $\|\mathbf{x}_i\| \to \infty$, the radial potential $\frac{\alpha}{2}\|\mathbf{x}_i - \boldsymbol{\mu}\|^2 \to \infty$ (quadratic growth), while the repulsion $-\gamma\lambda \sum_j \ln r_{ij}$ grows at most logarithmically. The erf attraction term is bounded by $\gamma \frac{\sigma\sqrt{\pi}}{2}\binom{n}{2}$. Therefore $E(\mathbf{X}) \to +\infty$ as any $\|\mathbf{x}_i\| \to \infty$, and E attains a finite infimum.

By the bounded monotone convergence theorem, the sequence $\{E(\mathbf{X}^t)\}$ converges. □

Remark. This does not imply convergence to a *global* minimum. The energy landscape has $O(n^2)$ pairwise interaction terms and is generically non-convex. Empirically, different initialisations converge to different local minima with Spearman $\rho \approx 0.65$, confirming the non-convexity.

C Force-Energy Verification: Why erf Is Necessary

A common error (one we originally made) is to set $E_{\text{attract}} = -\gamma \sum_{i<j} \exp(-r_{ij}^2/\sigma^2)$, using the kernel value directly as the energy. This is wrong:

$$\frac{d}{dr}\left[-\exp(-r^2/\sigma^2)\right] = \frac{2r}{\sigma^2}\exp(-r^2/\sigma^2) \neq \exp(-r^2/\sigma^2) \tag{29}$$

The gradient has an extra factor $2r/\sigma^2$. With this incorrect energy, the system is not gradient descent on any function, and energy can increase during iteration (observed: 14/14 steps energy *increasing*). The correct antiderivative is $\operatorname{erf}(r/\sigma)$:

$$\frac{d}{dr}\left[\frac{\sigma\sqrt{\pi}}{2}\operatorname{erf}(r/\sigma)\right] = \frac{\sigma\sqrt{\pi}}{2} \cdot \frac{2}{\sigma\sqrt{\pi}} e^{-r^2/\sigma^2} = e^{-r^2/\sigma^2} \quad \checkmark \tag{30}$$

After correcting the energy to use erf, energy decreases monotonically (14/14 steps down) with AUC unchanged ($\Delta = -0.00001$). The forces were correct all along. Only the energy accounting was wrong.